

Asynchronous programming

concurrency on one thread


What does a program do while it waits, and where is the rest of a function while it waits?

The model and its properties

one thread, four properties

The model

one thread, a stack, a heap, a queue, and a host



Two assumptions hold for the whole block.

One thread: all our code runs on it; no other thread runs code of ours.

An uncontrolled environment: the order of the host's events is not the program's to fix.

Suspension points

cooperative transfer of control

JAVASCRIPT

```
async function load() {  
  console.log("A");  
  const user = await fetchUser();  
  console.log("C", user);  
}
```

```
load();  
console.log("B");
```




the loop takes it, when the stack is empty

The rest, the highlighted line, is kept, registered on the awaited **promise**: a value that stands for a result to come.

Definition. An **asynchronous function** can be **suspended** at the points marked `await` in its text. At a **suspension point** its frame leaves the stack; **the part that remains to be run**, with its local variables, is kept in the heap. The caller resumes.

Property 1 · Suspension. An asynchronous function is suspended only at an `await` written in it. Its caller then goes on with its next statement: after `load();` the script prints B. A caller that writes `await load();` is suspended too.

Creating and awaiting a promise [†]`new Promise` and `await`



The loop runs the rest: `load()` resumes, and the `await` expression has `p`'s value.

`await e`: waiting and resuming

```
const user = await fetchUser();
```

JAVASCRIPT

1. It evaluates `e`. If the result is not a promise, it is wrapped into an already fulfilled one.
2. It registers the rest of the function on that promise, as one of its waiters: `await p` amounts to `p.then(rest)`^{*}.
3. It suspends the function: its frame leaves the stack, and its caller goes on.
4. When the promise is fulfilled, the rest goes to the microtask queue; the loop runs it, and `await e` has the promise's value.

[†] Python, C# and Rust have the same pair; what differs around it: slides 9 to 11, and the appendix.

^{*} `p.then(f)` registers `f`, to be called with the value of `p` once `p` is fulfilled.

Run-to-completion


the loop's rule, and two consequences

a and **b**: two requests, each answered after 300 ms; made in turn, or together before either is awaited.

in turn



together



Property 2 • Waiting is not running. Waits can overlap in time; executions of our code cannot: on one thread*, at any moment at most one computation is running.

Property 3 • Run-to-completion. The loop takes the next function from the queue only when the stack is empty: the running code has returned, and so has everything it called. Running code is never interrupted: it keeps the thread until it returns.

* JavaScript can run code of ours on another thread, a **worker**, with its own loop and its own heap.

Interleavings and atomicity

where one computation can run inside another

P's value arrives first



a suspension point:
the only place another
computation can run

Q's value arrives first



a stretch with no await:
atomic for every other
computation

■ running ▨ waiting

Execution: one sequence of resumptions. One program and one input have as many executions as the environment can produce.

Interleaving points are exactly the **suspension points**. Between two of them, nobody else runs.

Property 4 · Atomicity between suspension points. A stretch of a function's text with no `await` in it is **atomic** with respect to all other code of ours: nobody observes it half done. Others see all of it or none of it, and they can run between two stretches.

The race on one thread


a read, a suspension point, a write

```
JAVASCRIPT
let balance = 100;
async function deposit(n) {
  const b = balance;
  await approve(n);
  balance = b + n;
}
deposit(10);
deposit(20); // required: 130
```

`approve(n)` asks another service to approve the deposit; its answer comes 100 ms later.

Race condition: the outcome depends on an order the program does not fix. Here an **atomicity violation:** the deposit is split across a suspension point.

A **linearization point** of a call: an instant between the call and its return at which it takes effect. No choice of one for each deposit gives 120.



A lock? A thread's lock cannot help: there is one thread, and holding it across the wait blocks it.
An asynchronous lock, acquired with `await`, can.

Starting the work at the call, or later

Properties 1 to 4 describe JavaScript. Three questions, each decided by a small run, locate any other language with respect to them.

First: when does an asynchronous function start to run? At the call, or only when a scheduler takes it? The run: `load()`, then `console.log("B")`.

language	a call
JavaScript, C#	runs it up to its first <code>await</code> : A, B, C Ada
Python	a coroutine object, nothing runs: a warning, then B; <code>create_task</code> : B, A, C Ada
Rust	builds a future, which runs when polled
Kotlin	a <code>suspend</code> call waits: A, C Ada, B; <code>launch</code> queues it
Go, Java	a plain call waits: A, C Ada, B; with <code>go</code> or a virtual thread: B, A, C Ada

Two realizations of one model. In JavaScript and C# the call runs the function to its first `await`. In Python and Rust the call builds an object and runs nothing, session 5's thunk usable once: `await` runs it in place, `create_task` or `spawn` queues it.

The suspension point

at every `await`, or only when the value is missing

Second: what does an `await` do when the value is already there? Suspend anyway, so that every `await` is an interleaving point; or go straight on, so that the stretch stays atomic.

The run: the two deposits, with an `approve` that has nothing to wait for. In JavaScript, `await null` awaits a value already there; in Python, `asyncio.sleep(0)` asks for a wait of zero seconds.

language	an <code>await</code> on a finished result
JavaScript	suspends, even <code>await null</code> : 120
Python	goes straight on: 130; <code>await sleep(0)</code> suspends: 120
C#, Rust, Kotlin	goes straight on
Go, Java	there is no <code>await</code>

Of these languages only JavaScript suspends at every `await`. Elsewhere an `await` on a value already there is not a suspension point: on one thread, the stretch around it stays atomic.

The resuming thread

one thread, or a pool

Third: which thread resumes a suspended computation? The one thread of the loop, as the model assumes; or any thread of a pool.

The run: the two deposits, with a real wait, on the runtime each language gives by default.

language	the thread that resumes
JavaScript	the one thread: always 120
Python	the loop's one thread: always 120
C#	a thread of a pool: 110 or 120 in these runs
Rust, Kotlin	the executor's or the dispatcher's choice
Go, Java	several threads: 110 or 120 in these runs; 130 possible

Where a pool resumes the computations, two of them run at the same time: the model's first assumption is gone, and Property 4 with it, although every suspension point is written. Session 7.

Worked examples

the programs of part 1, run on the stage

1. One `await`
2. Two requests of 300 ms
3. A timer and a computation
4. Two deposits

Each one is run in the demo, step by step, on the model's parts: our code, the runtime, the host.

Example 1: one await

```
const fetchUser = () =>
  new Promise(res =>
    setTimeout(() => res("Ada"), 300));

async function load() {
  console.log("A");
  const user = await fetchUser();
  console.log("C", user);
}
await load();
console.log("B");
```

JavaScript ↗

Run

```
A
C Ada
B
```

`fetchUser` answers with `"Ada"` after 300 ms: a request to a server, imitated by a timer, which calls `res("Ada")` and so fulfills the promise.

What is printed, in which order?


1. A, B, C Ada
2. A, C Ada, B
3. B, A, C Ada

With `await load();` instead, the script is suspended too, and prints B last: A, C Ada, B.

Demo 1: the run of `load` step by step, on the model's four parts

`load` prints A, awaits a user that arrives after 300 ms, and prints C with the name. The script prints B.

Property 1: suspended only at the `await`; the caller continues. A, B, C Ada.



An `await` is this round trip: off the stack, into the heap, through the queue, back on the stack.



Example 2: two requests of 300 ms

```
async function inTurn() {  
  const a = await get("a");  
  const b = await get("b");  
  return [a, b];  
}  
  
async function together() {  
  const pa = get("a"), pb = get("b");  
  return [await pa, await pb];  
}  
  
let t = Date.now();  
console.log(await inTurn(), Date.now() - t);  
t = Date.now();  
console.log(await together(), Date.now() - t);
```

JavaScript ↗

`get` answers after 300 ms. In `together`, both requests have left before either is awaited.

How long does each function take?

1. 300 ms each
2. `inTurn` 600 ms, `together` 300 ms
3. 600 ms each


Run

```
["A", "B"] 600  
["A", "B"] 300
```

The text decides: a request written after an `await` leaves after it. Start the work first, then await.

Demo 2: two requests in turn, then together

Two values, a and b, each 300 ms away. inTurn awaits a, then asks for b; together asks for both first.



Property 2: the waits overlap, the running stretches never do. 600 against 300 ms.



The script prints 300 ms. The difference is where the calls stand relative to the await: start the work first, await afterwards.

Example 3: a timer and a computation

```
const t0 = Date.now();
setTimeout(() => console.log("tick at",
    Date.now() - t0, "ms"), 100);
// computes for 500 ms, never awaits
while (Date.now() - t0 < 500) {}
console.log("done at",
    Date.now() - t0, "ms");
```

JavaScript ↗

Run

recorded · node 25

```
done at 500 ms
tick at 504 ms
```

A timer due at 100 ms, then a loop that computes for 500 ms without reaching a suspension point.

When does `tick` print?

1. at 100 ms
2. at about 500 ms
3. never


Property 3: the timer's function is in the queue from 100 ms; the loop takes it only when the running code has returned. `done` at 500 ms, then `tick`.

Keeping the thread free is our job: split the computation, or move it to another thread, a worker.

Demo 3: a timer and a computation

A timer of 100 ms, then a loop that computes for 500 ms without an await.

Property 3: running code keeps the thread until it returns. `done` at 500 ms, then `tick`.



Cooperative: the loop runs the next function only where our code gives the thread back. Long work is split up, or moved to another thread.

Example 4: two deposits

```
let balance = 100;
async function deposit(n) {
  const b = balance;
  await approve(n);
  balance = b + n;
}
deposit(10);
deposit(20);
```

JavaScript ↗

Run

Run, first approval slower

balance 110

`approve` answers after 100 ms. Required: 130 when both calls have finished.

What is the balance then?

1. 130
2. 120
3. 110


It happens wherever we read shared state, wait, then write it:
a counter updated after a request, a cache refreshed after a fetch.

Demo 4: two deposits

equal approvals; then the first one slower

Two deposits on one balance, 10 and 20, each suspended at its approval. Required: 130.

Property 4 ends at the **await**: the read and the write are two stretches, and the other call runs between them. 120.



120, where 130 was required: each call was suspended between its read and its write, and the other ran there. First approval slower: 110.

The same programs in Python

a bare call; an `await` on a finished result

```
async def main():
    load()                # the call alone
    print("B")
    await asyncio.sleep(0.5)
    asyncio.create_task(load()) # queued
    print("B")
    await asyncio.sleep(0.5)
```

Python ↗

Run

```
<exec>:10: RuntimeWarning: coroutine 'load' was never awaited
RuntimeWarning: Enable tracemalloc to get the object allocation
B
B
A
C Ada
```

Does a call start the work?

A coroutine object that is never awaited never runs: its work is lost, and the warning is the only sign of it.

```
async def trial(name, approve):
    account = {"balance": 100}
    async def deposit(n):
        b = account["balance"]
        await approve()
        account["balance"] = b + n
    await asyncio.gather(
        deposit(10), deposit(20))
    print(name, account["balance"])
```

Python ↗

Run

```
await trial("nothing to wait for:", done)
await trial("waits 0 seconds:", sleep0)
```

```
nothing to wait for: 130
waits 0 seconds: 120
```

Is every `await` a suspension point?

Conclusion

the four properties; what is left open; the lab, and session 7

Conclusion of block 2

- 1 Suspended only at an `await` in it; after `load()`; the caller goes on, after `await load()`; it waits.
- 2 Waits overlap; on one thread, running never does.
- 3 Run-to-completion: running code keeps the thread until it returns.
- 4 A stretch with no `await` is atomic. The race on one thread is an atomicity violation, not a data race.

the language guarantees

atomicity between suspension points, given one thread

a library gives

combinators over promises:
`Promise.all`, `gather`

the programmer keeps

every invariant that must survive a wait

Left open: what the kept rest of a function is; who holds it while the function waits; who invokes it, and when. The lab's optional quest.

Next

the lab, and session 7

- lab** The warm-up: the two blocks of today, on JavaScript with Python beside it. The optional quest: a small page on continuations, with exercises on paper; the appendix of this deck is its reading.
-
- 7** Several threads on one memory. The repaired deposit, run by two Java threads twenty thousand times: 130 in about 98 runs of 100, 110 or 120 in the others. The suspension points are gone, and with them the atomic stretch.

The suspended computation as a value

continuations, a promise, a loop, a generator and a driver

Continuations

the kept rest, written by hand

```
function load() {  
  console.log("A");  
  return fetchUser().then(user => { console.log("C", user); });  
}  
load();  
console.log("B");
```

JavaScript ↗

Run

A
B
C Ada

Definition. The **continuation** of a point in a computation is the rest of the computation from that point, as a function of the value expected there.

In `double(x) + 1`, while `double(x)` is computed, the rest is the context `[·] + 1`: the function `v => v + 1`. At the `await` of `load`, the rest is `user => { console.log("C", user); }`.

Continuation-passing style

a function that cannot return takes its continuation

JAVASCRIPT

```
// doubleK asks a server for 2x; the answer goes to k, 100 ms later  
const doubleK = (x, k) => setTimeout(() => k(2 * x), 100);  
const f = x => double(x) + 1; // double cannot return: write fK
```

```
const fK = (x, k) => doubleK(x, v => k(v + 1)); // A  
const fK = (x, k) => doubleK(x, k) + 1;        // B  
const fK = (x, k) => k(doubleK(x) + 1);        // C
```

JavaScript ↗

Run

```
recorded · node 25  A gives k the value 7  
B gives k the value 6  
C gives k the value 11, then fails: k is not a function
```

The cost: every function that needs the result of `fK` for its own result takes a continuation too, up the chain of calls. Hence `async` on every function that uses `await`, and on every caller that awaits its result.

A promise by hand

a write-once cell that holds continuations

```
class P {  
  state = "pending"; value; waiters = [];  
  then(k) {  
    if (this.state === "pending") this.waiters.push(k);  
    else later(() => k(this.value));  
  }  
  resolve(v) {  
    if (this.state !== "pending") return;  
    this.state = "fulfilled"; this.value = v;  
    for (const k of this.waiters) later(() => k(v));  
  }  
}
```

JavaScript ↗

A promise nobody resolves: its waiters never run.

If nothing can reach it any more, the collector frees it with them, silently. If something still can, they stay for ever: a memory leak.

`resolve` has the waiting continuations. Call each of them now, inside `resolve`, or put each in the queue for the loop?

The queue and the loop the scheduler

JavaScript ↗

```
const ready = []; let scheduled = false;
function later(k) {                      // put k in the queue
  ready.push(k);
  if (!scheduled) { scheduled = true; setTimeout(run, 0); }
}
function run() {                        // the loop: one at a time, to its end
  scheduled = false;
  while (ready.length > 0) { const k = ready.shift(); k(); }
}
```

Run the test, both ways

recorded · node 25

```
call now   : k 1 | after resolve
queue      : after resolve | k 1
```

JavaScript queues: a function given to `then` is never called before the current code has run to completion (Promises/A+, 2.2.4). Property 3 is these lines. Property 4 is kept by `then` and `resolve` calling `later` instead of `k`.

A value sent into a generator

session 5's object, resumed with a value

```
function* two() {  
  const x = yield 1;  
  console.log("x is", x);  
}  
const g = two();  
console.log(g.next());  
console.log(g.next("Ada"));
```

JavaScript ↗

Run

```
{ value: 1, done: false }  
x is Ada  
{ value: undefined, done: true }
```

`g.next(v)` makes `v` the value of the `yield` where `g` was suspended.

```
function* loads() {  
  for (let i = 0; i < 3; i++) {  
    const user = yield fetchUser();  
    console.log(i, user);  
  }  
}  
drive(loads);  
console.log("B");
```

JavaScript ↗

Run

```
B  
0 Ada  
1 Ada  
2 Ada
```

A wait inside a loop: the generator's object keeps `i` and the position.

An async function as a generator and a driver

await became yield

```
function drive(gen) {
  const g = gen();
  const step = v => {
    const r = g.next(v);
    if (!r.done) r.value.then(step);
  };
  step();
}
```

JavaScript ↗

Run

A
B
C Ada

code

VE


```
function drive(gen) {
  const g = gen();
  const step = v => {
    const r = g.next(v);
    if (!r.done) r.value.then(step);
  };
  step();
}
```

IPT

```
function* load() {
  console.log("A");
  const user = yield fetchUser();
  console.log("C", user);
}
```

SOLE

A



A, B, C Ada on our promise, our queue, our loop and our driver: what the language did.

run, our loop, takes step from ready and calls it. g.next("Ada") puts Ada where the yield was, and load prints C Ada.

Languages against the model

task management and stack management

language	a call	an <code>await</code> on a finished result	the thread that resumes	who keeps the suspended state
JavaScript	runs to the first <code>await</code>	suspends	the one thread	the compiler (stackless)
C#		goes straight on	a thread of a pool	
Python (asyncio)			the loop's one thread	
Rust			the executor's choice	
Kotlin	a <code>suspend</code> call waits; <code>launch</code> queues		the dispatcher's choice	
Go, Java (virtual threads)	a call waits; <code>go</code> or a virtual thread starts it apart	no <code>await</code>	several threads	the runtime (stackful)
OCaml 5 (effects)	runs until it performs an effect		one domain; a scheduler in a library (not run)	

Task management: who may run between two of our lines. **Stack management:** who keeps the suspended state. Two independent questions (Adya et al., 2002).

Worker threads

another thread, another loop

main.mjs ↗

```
import { Worker } from "node:worker_threads";
const t0 = Date.now(), ms = () => Date.now() - t0;
const w = new Worker(
  new URL("./fib.mjs", import.meta.url));
w.on("message", () => {
  console.log("fib(40) at", ms(), "ms");
  w.terminate();
});
w.postMessage(40);
setTimeout(() =>
  console.log("tick at", ms(), "ms"), 100);
```

fib.mjs ↗

```
import { parentPort } from "node:worker_threads";
const fib = n =>
  n < 2 ? n : fib(n - 1) + fib(n - 2);
parentPort.on("message", n =>
  parentPort.postMessage(fib(n)));
```

A **worker** is another thread with its own loop: its own stack, heap and queue.

The two threads share no object. `postMessage` sends a **copy**; it arrives as a function in the other thread's queue.

Run

recorded · node 25

tick at 104 ms

fib(40) at 563 ms

The tick is on time: the computation runs on the worker's thread. On one thread it came at 504 ms.

On each thread, Properties 1 to 4 hold; between threads, messages: session 8. The exception is a `SharedArrayBuffer`, memory both threads write; with it, data races return: session 7.