

# Asynchronous programming

concurrency on one thread

---


What does a program do while it waits, and where is the rest of a function while it waits?

# The model and its properties

one thread, four properties

# The model

one thread, a stack, a heap, a queue, and a host



Two assumptions hold for the whole block.

**One thread:** all our code runs on it; no other thread runs code of ours.

**An uncontrolled environment:** the order of the host's events is not the program's to fix.


# Suspension points

cooperative transfer of control

JAVASCRIPT

```
async function load() {  
  console.log("A");  
  const user = await fetchUser();  
  console.log("C", user);  
}
```

```
load();  
console.log("B");
```




the loop takes it, when the stack is empty

The rest, the highlighted line, is kept, registered on the awaited **promise**: a value that stands for a result to come.

**Definition.** An **asynchronous function** can be **suspended** at the points marked `await` in its text. At a **suspension point** its frame leaves the stack; **the part that remains to be run**, with its local variables, is kept in the heap. The caller resumes.

**Property 1 · Suspension.** An asynchronous function is suspended only at an `await` written in it. Its caller then goes on with its next statement: after `load();` the script prints B. A caller that writes `await load();` is suspended too.

# Creating and awaiting a promise <sup>†</sup>`new Promise` and `await`



The loop runs the rest: `load()` resumes, and the `await` expression has `p`'s value.

`await e`: waiting and resuming

```
const user = await fetchUser();
```

JAVASCRIPT

1. It evaluates `e`. If the result is not a promise, it is wrapped into an already fulfilled one.
2. It registers the rest of the function on that promise, as one of its waiters: `await p` amounts to `p.then(rest)`<sup>\*</sup>.
3. It suspends the function: its frame leaves the stack, and its caller goes on.
4. When the promise is fulfilled, the rest goes to the microtask queue; the loop runs it, and `await e` has the promise's value.

<sup>†</sup> Python, C# and Rust have the same pair; what differs around it: part 3, and the appendix.

<sup>\*</sup> `p.then(f)` registers `f`, to be called with the value of `p` once `p` is fulfilled.

# Callbacks

the rest of a computation, passed by hand

Before promises, a function whose result comes later took a **callback**: the rest of the computation, as a function, which it calls with the result.

1. **Nesting**: each step is written inside the callback of the step before. With `await`, the steps follow one another.
2. **No composition**: a result not yet arrived is not a value; waiting for two takes a counter. A promise is a value: stored, passed on, awaited by several.
3. **Inversion of control**: the function we call decides when our callback runs, and how many times. A promise is settled once, and what waits on it always goes to the queue.
4. **Errors**: an exception thrown in a callback does not reach the caller's `try`. Failure: session 9.

```
try {  
  setTimeout(() => {  
    throw new Error("no user");  
  }, 100);  
} catch (e) {  
  console.log("caught", e.message);  
}  
console.log("B");
```

JavaScript ↗

Run

```
recorded · node 25  B  
Error: no user  
    at Timeout._onTimeout (callbacks.mjs:39:15)
```

# Run-to-completion


the loop's rule, and two consequences

**a** and **b**: two requests, each answered after 300 ms; made in turn, or together before either is awaited.

**in turn**



**together**



**Property 2 • Waiting is not running.** Waits can overlap in time; executions of our code cannot: on one thread\*, at any moment at most one computation is running.

**Property 3 • Run-to-completion.** The loop takes the next function from the queue only when the stack is empty: the running code has returned, and so has everything it called. Running code is never interrupted: it keeps the thread until it returns.

\* JavaScript can run code of ours on another thread, a **worker**, with its own loop and its own heap.

# Interleavings and atomicity

where one computation can run inside another

## P's value arrives first



a suspension point:  
the only place another  
computation can run

## Q's value arrives first



a stretch with no await:  
atomic for every other  
computation

■ running    ▨ waiting

**Execution:** one sequence of resumptions. One program and one input have as many executions as the environment can produce.

**Interleaving points** are exactly the **suspension points**. Between two of them, nobody else runs.

**Property 4 · Atomicity between suspension points.** A stretch of a function's text with no `await` in it is **atomic** with respect to all other code of ours: nobody observes it half done, except the functions it calls, which run inside it. Others see all of it or none of it, and they can run between two stretches.

# The race on one thread


a read, a suspension point, a write

```
JAVASCRIPT
let balance = 100;
async function deposit(n) {
  const b = balance;
  await approve(n);
  balance = b + n;
}
deposit(10);
deposit(20); // required: 130
```

`approve(n)` asks another service to approve the deposit; its answer comes 100 ms later.

**Race condition:** the outcome depends on an order the program does not fix. Here an **atomicity violation:** the deposit is split across a suspension point.

A **linearization point** of a call: an instant between the call and its return at which it takes effect. No choice of one for each deposit gives 120.



A lock? A thread's lock cannot help: there is one thread, and holding it across the wait blocks it.  
An asynchronous lock, acquired with `await`, can.

# Worked examples

the programs of part 1, run on the stage

1. One `await`
2. Two requests of 300 ms
3. A timer and a computation
4. Two deposits

Each one is run in the demo, step by step, on the model's parts: our code, the runtime, the host.

# Example 1: one await

```
const fetchUser = () =>
  new Promise(res =>
    setTimeout(() => res("Ada"), 300));

async function load() {
  console.log("A");
  const user = await fetchUser();
  console.log("C", user);
}
await load();
console.log("B");
```

JavaScript ↗

Run

```
A
C Ada
B
```

`fetchUser` answers with `"Ada"` after 300 ms: a request to a server, imitated by a timer, which calls `res("Ada")` and so fulfills the promise.

What is printed, in which order?


1. A, B, C Ada
2. A, C Ada, B
3. B, A, C Ada

With `await load();` instead, the script is suspended too, and prints B last: A, C Ada, B.

# Demo 1: the run of `load` step by step, on the model's four parts

`load` prints A, awaits a user that arrives after 300 ms, and prints C with the name. The script prints B.

Property 1: suspended only at the `await`; the caller continues. A, B, C Ada.



An `await` is this round trip: off the stack, into the heap, through the queue, back on the stack.



## Example 2: two requests of 300 ms

```
async function inTurn() {  
  const a = await get("a");  
  const b = await get("b");  
  return [a, b];  
}  
  
async function together() {  
  const pa = get("a"), pb = get("b");  
  return [await pa, await pb];  
}  
  
let t = Date.now();  
console.log(await inTurn(), Date.now() - t);  
t = Date.now();  
console.log(await together(), Date.now() - t);
```

JavaScript ↗

`get` answers after 300 ms. In `together`, both requests have left before either is awaited.

How long does each function take?

1. 300 ms each
2. `inTurn` 600 ms, `together` 300 ms
3. 600 ms each


Run

```
["A", "B"] 600  
["A", "B"] 300
```

The text decides: a request written after an `await` leaves after it. Start the work first, then await.

# Demo 2: two requests in turn, then together

Two values, a and b, each 300 ms away. inTurn awaits a, then asks for b; together asks for both first.



Property 2: the waits overlap, the running stretches never do. 600 against 300 ms.

The script prints 300 ms. The difference is where the calls stand relative to the await: start the work first, await afterwards.

# Example 3: a timer and a computation

```
const t0 = Date.now();
setTimeout(() => console.log("tick at",
    Date.now() - t0, "ms"), 100);
// computes for 500 ms, never awaits
while (Date.now() - t0 < 500) {}
console.log("done at",
    Date.now() - t0, "ms");
```

JavaScript ↗

Run

recorded · node 25

```
done at 500 ms
tick at 504 ms
```

A timer due at 100 ms, then a loop that computes for 500 ms without reaching a suspension point.

When does `tick` print?

1. at 100 ms
- 2. at about 500 ms**
3. never


Property 3: the timer's function is in the queue from 100 ms; the loop takes it only when the running code has returned. `done` at 500 ms, then `tick`.

Keeping the thread free is our job: split the computation, or move it to another thread, a worker.

# Demo 3: a timer and a computation

A timer of 100 ms, then a loop that computes for 500 ms without an await.

Property 3: running code keeps the thread until it returns. `done` at 500 ms, then `tick`.



Cooperative: the loop runs the next function only where our code gives the thread back. Long work is split up, or moved to another thread.

# Example 4: two deposits

```
let balance = 100;
async function deposit(n) {
  const b = balance;
  await approve(n);
  balance = b + n;
}
deposit(10);
deposit(20);
```

JavaScript ↗

Run

Run, first approval slower

balance 110

`approve` answers after 100 ms. Required: 130 when both calls have finished.

What is the balance then?

1. 130
2. 120
3. 110


It happens wherever we read shared state, wait, then write it:  
a counter updated after a request, a cache refreshed after a fetch.

# Demo 4: two deposits

equal approvals; then the first one slower

Two deposits on one balance, 10 and 20, each suspended at its approval. Required: 130.

Property 4 ends at the **await**: the read and the write are two stretches, and the other call runs between them. 120.



120, where 130 was required: each call was suspended between its read and its write, and the other ran there. First approval slower: 110.

# Other languages

the same programs, against the model

# The same programs in Python

a bare call; an `await` on a finished result

```
async def main():
    load()                # the call alone
    print("B")
    await asyncio.sleep(0.5)
    asyncio.create_task(load()) # queued
    print("B")
    await asyncio.sleep(0.5)
```

Python ↗

Run

```
<exec>:10: RuntimeWarning: coroutine 'load' was never awaited
RuntimeWarning: Enable tracemalloc to get the object allocation
B
B
A
C Ada
```

Does a call start the work?

A coroutine object that is never awaited never runs: its work is lost, and the warning is the only sign of it.

```
async def trial(name, approve):
    account = {"balance": 100}
    async def deposit(n):
        b = account["balance"]
        await approve()
        account["balance"] = b + n
    await asyncio.gather(
        deposit(10), deposit(20))
    print(name, account["balance"])
```

Python ↗

Run

```
await trial("nothing to wait for:", done)
await trial("waits 0 seconds:", sleep0)
```

```
nothing to wait for: 130
waits 0 seconds: 120
```

Is every `await` a suspension point?

# Starting the work at the call, or later

Properties 1 to 4 describe JavaScript. Four questions, each decided by a small run or by the language's interface, locate any other language with respect to them.

First: when does an asynchronous function start to run? At the call, or only when a scheduler takes it? The run: `load()`, then `console.log("B")`.

| language       | a call                                                                                                                          |
|----------------|---------------------------------------------------------------------------------------------------------------------------------|
| JavaScript, C# | runs it up to its first <code>await</code> : A, B, C Ada                                                                        |
| Python         | a coroutine object, nothing runs: a warning, then B; <code>create_task</code> : B, A, C Ada                                     |
| Rust           | builds a future, which runs when polled                                                                                         |
| Kotlin         | a <code>suspend</code> call waits: A, C Ada, B; <code>launch</code> queues it                                                   |
| Go, Java       | a plain call waits: A, C Ada, B; with <code>go</code> or a virtual thread: B, A, C Ada in these runs, A and B in no fixed order |

Two realizations of one model. In JavaScript and C# the call runs the function to its first `await`. In Python and Rust the call builds an object and runs nothing, session 5's thunk usable once: `await` runs it in place, `create_task` or `spawn` queues it.

# The suspension point

at every `await`, or only when the value is missing

Second: what does an `await` do when the value is already there? Suspend anyway, so that every `await` is an interleaving point; or go straight on, so that the stretch stays atomic.

The run: the two deposits, with an `approve` that has nothing to wait for. In JavaScript, `await null` awaits a value already there; in Python, `asyncio.sleep(0)` asks for a wait of zero seconds.

| language         | an <code>await</code> on a finished result                       |
|------------------|------------------------------------------------------------------|
| JavaScript       | suspends, even <code>await null</code> : 120                     |
| Python           | goes straight on: 130; <code>await sleep(0)</code> suspends: 120 |
| C#, Rust, Kotlin | goes straight on                                                 |
| Go, Java         | there is no <code>await</code>                                   |

Of these languages only JavaScript suspends at every `await`. Elsewhere an `await` on a value already there is not a suspension point: on one thread, the stretch around it stays atomic.

# The resuming thread

one thread, or a pool

Third: which thread resumes a suspended computation? The one thread of the loop, as the model assumes; or any thread of a pool.

The run: the two deposits, with a real wait, on the runtime each language gives by default.

| language     | the thread that resumes                                            |
|--------------|--------------------------------------------------------------------|
| JavaScript   | the one thread: always 120                                         |
| Python       | the loop's one thread: always 120                                  |
| C#           | in a console program, a thread of a pool: 110 or 120 in these runs |
| Rust, Kotlin | the executor's or the dispatcher's choice                          |
| Go, Java     | several threads: 110 or 120 in these runs; 130 possible            |

Where a pool resumes the computations, two of them run at the same time: the model's first assumption is gone, and Property 4 with it, although every suspension point is written. Session 7.

# Registering the rest

then, and what other languages have

Fourth: how is the rest of a function handed to the result it waits for? A call registers it on the result, as `p.then(k)` does; or a scheduler asks the result whether it is ready.

Before `await`, this was all we had. Each language reached it from a different start.

| language   | the rest, handed over                                |
|------------|------------------------------------------------------|
| JavaScript | <code>then</code> ; before 2015, a callback          |
| C#         | <code>ContinueWith</code> ; <code>Wait</code> blocks |
| Java       | <code>thenApply</code> ; <code>get</code> blocks     |
| Python     | <code>add_done_callback</code>                       |
| Kotlin     | a <code>Continuation</code> , from the compiler      |
| Rust       | none: the executor polls the future                  |

JavaScript could not block its one thread: passing the rest was the only way to wait. Languages with threads could block first, and added this later. With `await`, the compiler writes the registration for us.

# Conclusion

the four properties; what is left open; the lab, and session 7

# Conclusion of block 2

- 1 Suspended only at an `await` in it; after `load()`; the caller goes on, after `await load()`; it waits.
- 2 Waits overlap; on one thread, running never does.
- 3 Run-to-completion: running code keeps the thread until it returns.
- 4 A stretch with no `await` is atomic. The race on one thread is an atomicity violation, not a data race.

## the language guarantees

atomicity between suspension points, given one thread

## a library gives

combinators over promises:  
`Promise.all`, `gather`

## the programmer keeps

every invariant that must survive a wait

**Left open:** what the kept rest of a function is; who holds it while the function waits; who invokes it, and when. The lab's optional quest.

# Next

the lab, and session 7

- lab** The warm-up: the two blocks of today, on JavaScript with Python beside it. The optional quest: a small page on continuations, with exercises on paper; the appendix of this deck is its reading.
- 
- 7** Several threads on one memory. The repaired deposit, run by two Java threads twenty thousand times: 130 in about 98 runs of 100, 110 or 120 in the others. The suspension points are gone, and with them the atomic stretch.

# The suspended computation as a value

continuations, a promise, a loop, a generator and a driver

# Continuations

the kept rest, written by hand

```
function load() {  
  console.log("A");  
  return fetchUser().then(user => { console.log("C", user); });  
}  
load();  
console.log("B");
```

JavaScript ↗

Run

A  
B  
C Ada

**Definition.** The **continuation** of a point in a computation is the rest of the computation from that point, as a function of the value expected there.

In `double(x) + 1`, while `double(x)` is computed, the rest is the context `[·] + 1`: the function `v => v + 1`. At the `await` of `load`, the rest is `user => { console.log("C", user); }`.

# Continuation-passing style

a function that cannot return takes its continuation

JAVASCRIPT

*// doubleK asks a server for 2x; the answer goes to k, 100 ms later*

```
const doubleK = (x, k) => setTimeout(() => k(2 * x), 100);
```

```
const f = x => double(x) + 1;    // double cannot return: write fK
```

```
const fK = (x, k) => doubleK(x, v => k(v + 1));    // A
```

```
const fK = (x, k) => doubleK(x, k) + 1;           // B
```

```
const fK = (x, k) => k(doubleK(x) + 1);           // C
```

JavaScript ↗

Run

recorded · node 25   A gives k the value 7

B gives k the value 6

C gives k the value 11, then fails: k is not a function

**The cost:** every function that needs the result of `fK` for its own result takes a continuation too, up the chain of calls. Hence `async` on every function that uses `await`, and on every caller that awaits its result.

# A promise by hand

a write-once cell that holds continuations

```
class P {  
  state = "pending"; value; waiters = [];  
  then(k) {  
    if (this.state === "pending") this.waiters.push(k);  
    else later(() => k(this.value));  
  }  
  resolve(v) {  
    if (this.state !== "pending") return;  
    this.state = "fulfilled"; this.value = v;  
    for (const k of this.waiters) later(() => k(v));  
  }  
}
```

JavaScript ↗

A promise nobody resolves: its waiters never run.

If nothing can reach it any more, the collector frees it with them, silently. If something still can, they stay for ever: a memory leak.

`resolve` has the waiting continuations. Call each of them now, inside `resolve`, or put each in the queue for the loop?

# The queue and the loop the scheduler

JavaScript ↗

```
const ready = []; let scheduled = false;
function later(k) { // put k in the queue
  ready.push(k);
  if (!scheduled) { scheduled = true; setTimeout(run, 0); }
}
function run() { // the loop: one at a time, to its end
  scheduled = false;
  while (ready.length > 0) { const k = ready.shift(); k(); }
}
```

Run the test, both ways

recorded · node 25

```
call now   : k 1 | after resolve
queue      : after resolve | k 1
```

JavaScript queues: a function given to `then` is never called before the current code has run to completion (Promises/A+, 2.2.4). Property 3 is these lines. Property 4 is kept by `then` and `resolve` calling `later` instead of `k`.

# A value sent into a generator

session 5's object, resumed with a value

```
function* two() {  
  const x = yield 1;  
  console.log("x is", x);  
}  
const g = two();  
console.log(g.next());  
console.log(g.next("Ada"));
```

JavaScript ↗

Run

```
{ value: 1, done: false }  
x is Ada  
{ value: undefined, done: true }
```

`g.next(v)` makes `v` the value of the `yield` where `g` was suspended.

```
function* loads() {  
  for (let i = 0; i < 3; i++) {  
    const user = yield fetchUser();  
    console.log(i, user);  
  }  
}  
drive(loads);  
console.log("B");
```

JavaScript ↗

Run

```
B  
0 Ada  
1 Ada  
2 Ada
```

A wait inside a loop: the generator's object keeps `i` and the position.

# An async function as a generator and a driver

await became yield

```
function drive(gen) {
  const g = gen();
  const step = v => {
    const r = g.next(v);
    if (!r.done) r.value.then(step);
  };
  step();
}
```

JavaScript ↗

Run

A  
B  
C Ada

code

VE


```
function drive(gen) {
  const g = gen();
  const step = v => {
    const r = g.next(v);
    if (!r.done) r.value.then(step);
  };
  step();
}
```

IPT

```
function* load() {
  console.log("A");
  const user = yield fetchUser();
  console.log("C", user);
}
```

SOLE

A



A, B, C Ada on our promise, our queue, our loop and our driver: what the language did.

run, our loop, takes step from ready and calls it. g.next("Ada") puts Ada where the yield was, and load prints C Ada.

# Languages against the model

task management and stack management

| language                   | a call                                                            | an <code>await</code> on a finished result | the thread that resumes                        | who keeps the suspended state |
|----------------------------|-------------------------------------------------------------------|--------------------------------------------|------------------------------------------------|-------------------------------|
| JavaScript                 | runs to the first <code>await</code>                              | suspends                                   | the one thread                                 | the compiler (stackless)      |
| C#                         |                                                                   | goes straight on                           | a thread of a pool (console)                   |                               |
| Python (asyncio)           |                                                                   |                                            | the loop's one thread                          |                               |
| Rust                       |                                                                   |                                            | the executor's choice                          |                               |
| Kotlin                     | a <code>suspend</code> call waits; <code>launch</code> queues     |                                            | the dispatcher's choice                        |                               |
| Go, Java (virtual threads) | a call waits; <code>go</code> or a virtual thread starts it apart | no <code>await</code>                      | several threads                                | the runtime (stackful)        |
| OCaml 5 (effects)          | runs until it performs an effect                                  |                                            | one domain; a scheduler in a library (not run) |                               |

**Task management:** who may run between two of our lines. **Stack management:** who keeps the suspended state. Two independent questions (Adya et al., 2002).

# Worker threads

another thread, another loop

main.mjs ↗

```
import { Worker } from "node:worker_threads";
const t0 = Date.now(), ms = () => Date.now() - t0;
const w = new Worker(
  new URL("./fib.mjs", import.meta.url));
w.on("message", () => {
  console.log("fib(40) at", ms(), "ms");
  w.terminate();
});
w.postMessage(40);
setTimeout(() =>
  console.log("tick at", ms(), "ms"), 100);
```

fib.mjs ↗

```
import { parentPort } from "node:worker_threads";
const fib = n =>
  n < 2 ? n : fib(n - 1) + fib(n - 2);
parentPort.on("message", n =>
  parentPort.postMessage(fib(n)));
```

A **worker** is another thread with its own loop: its own stack, heap and queue.

The two threads share no object. `postMessage` sends a **copy**; it arrives as a function in the other thread's queue.

Run

recorded · node 25

tick at 104 ms

fib(40) at 563 ms

The tick is on time: the computation runs on the worker's thread. On one thread it came at 504 ms.

On each thread, Properties 1 to 4 hold; between threads, messages: session 8. The exception is a `SharedArrayBuffer`, memory both threads write; with it, data races return: session 7.

# The repair

the read after the wait

```
async function deposit(n) {  
  const b = balance;  
  await approve(n);  
  balance = b + n;  
}  
async function repaired(n) {  
  await approve(n);  
  const b = balance;  
  balance = b + n;  
}
```

JavaScript ↗

Run

```
deposit first approval 100 ms: 120  
deposit first approval 200 ms: 110  
repaired first approval 100 ms: 130  
repaired first approval 200 ms: 130
```

The read moves after the `await`. The read and the write are then one stretch with no `await`: atomic, by Property 4.

Each deposit takes effect in that stretch. 130, whichever approval arrives first.

What was assumed before an `await` is re-established after it.

# Inversion of control

what a promise rules out

the callee calls it

with a callback

with a promise

twice

the rest runs twice

settled once: a second `res` does nothing

now or later

the order of our lines depends on it

what waits always goes to the queue

never

it waits for ever, silently

the same

```
const user = id =>
  new Promise(res => userK(id, res));
const show = () => {
  user(1).then(u => console.log("C", u));
  console.log("B");
};
show();
setTimeout(show, 200);
```

JavaScript ↗

Run

```
B
C Ada
B
C Ada
```

The cached `userK` of the callbacks slide, wrapped in a promise: B first, on a miss and on a hit.

The first two are ruled out by the promise's two rules. The third is a liveness failure: no finite run shows it.