

Warm-up: an email as a type, then five small types

Programming Paradigms
Session 3 — the warm-up, before the quest

Part 1 on paper, in pairs; Part 2 at the keyboard, with `warmup.ml`.

SOLUTIONS

The lecture said that an algebraic data type is a closed list of shapes, written where the compiler can read it, and that the compiler pays us back by checking every `match`. Before the music language, we try that on something everyone owns: a notification that goes out by email, by SMS or by push. Nothing here is themed, and in Part 1 there is nothing to run. The notation is OCaml's, and these are all we need:

<code>type t = A B C</code>	a <i>sum</i> : a value of <code>t</code> is A, or B, or C
<code>type t = A of int B of string</code>	a sum whose shapes carry something
<code>type t = { f : int; g : bool }</code>	a <i>product</i> : one value holds an <code>f</code> and a <code>g</code>
<code>int * bool</code>	a product without names, a pair
<code>t list, t option</code>	zero or more <code>t</code> ; a <code>t</code> or nothing (<code>Some x / None</code>)

Part 1 — on paper

Exercise 1 (the message)

Write OCaml types for the following, choosing sum or product as the facts demand. An **address** is a user name and a domain, both text. A **channel** is one of three: email, to an address; SMS, to a phone number; push, to a device token (both text). A **message** has a channel, a body, and a subject — which an SMS does not have. Then draw one concrete message sent by email as a tree: constructors at the nodes, text at the leaves. (If you write `type phone = string`, note that this *synonym* invents no values: every string is a phone. A sum does invent values — `Email`, `Sms`, `Push` did not exist before.)

Answer.

```
type address = { user : string; domain : string }
type channel = Email of address | Sms of string | Push of string
type message = { to_ : channel; subject : string option; body : string }
```

Address and message are products (every field is always present); channel is a sum (exactly one of three kinds, each carrying its own kind of destination — the payloads differ, which is why a sum and not a record with a “kind” field). The subject is `string option`, not an empty string: the type says it can be absent, and every function that reads it has to face that case. The tree has `message` at the root; below it `Email` with an `address` node holding two text leaves, `Some "..."` or `None`, and the body text.

Exercise 2 (why “algebraic”)

Let

```
type priority = Low | Normal | High
```

How many distinct values are there of each type: `bool * bool`; `priority`; `priority * bool`; `priority option`; `bool option option`; `(priority * bool) option`? State the two rules that give the answers.

Answer. 4; 3; 6; 4; 3; 7. A *sum* has as many values as its summands added: `option` adds one (`None`) to what it wraps, so `priority option` has $3 + 1$ and `bool option option` has $(2 + 1) + 1$.

A *product* multiplies: `priority * bool` has 3×2 , and wrapping it in `option` gives $6 + 1$. That is the algebra in the name: a bare constructor counts 1, `|` is $+$, $*$ and record fields are $\times$. (A recursive type, or one containing `string`, has infinitely many values — the rules still say how each is built.)

Exercise 3 (missing and unreachable)

Two functions on `channel`, written in a hurry:

```
let destination c = match c with
| Email a -> a.user ^ "@" ^ a.domain
| Sms number -> number

let is_push c = match c with
| Email _ | Sms _ | Push _ -> false
| Push _ -> true
```

For each: does it compile? What does the compiler print, and what is the message called? What does `destination (Push t)` do at run time, and what does `is_push (Push t)` return? Fix both.

Answer. Both compile — with a warning each, and the two warnings are the two dual checks of the lecture. The first is *Warning 8: this pattern-matching is not exhaustive. Here is an example of a case that is not matched: Push _*; at run time `destination (Push t)` raises `Match_failure`. The second is *Warning 11: this match case is unused*, on the arm `Push _ -> true`: the or-pattern above already takes every `Push`, so `is_push (Push t)` returns `false` — no crash, only a wrong answer, and only the warning reveals it. Fixes: add `| Push token -> token`; and take `Push _` out of the or-pattern (or move its arm above it). Arms are tried top to bottom, and the first that fits wins.

Exercise 4 (the grid)

Three channels, `Email`, `Sms`, `Push`, and two things to do with each: `send` and `format`. Six cells of code. Draw the table, channels as rows and operations as columns. In OCaml, which cells sit together in one function? In Java, with one class per channel, which sit together in one class? Now the boss asks for two changes, separately: *add WhatsApp as a channel*, and *add schedule as an operation*. For each change, in each language: how many new cells, how many *existing* functions or classes are touched — and what does the compiler say if a cell is forgotten?

Answer. The grid, with the two changes marked:

	send	format	<i>schedule</i> (new)
Email	●	●	○
Sms	●	●	○
Push	●	●	○
<i>WhatsApp</i> (new)	○	○	○

- the six cells that exist; ○ what a change adds. The two layouts, as code:

OCaml — a column is one function

```
let send c body = match c with
| Email a -> smtp a body
| Sms n -> sms_gateway n body
| Push t -> push_service t body

let format c body = match c with
| Email _ -> body ^ "\n-- sent by mail"
| Sms _ -> String.sub body 0 160
| Push _ -> body
```

Java — a row is one class

```
interface Channel {
    void send(String body);
    String format(String body);
}

class Email implements Channel {
    Address a;
    public void send(String body) { smtp(a, body); }
    public String format(String body)
        { return body + "\n-- sent by mail"; }
}

class Sms implements Channel { ... }
class Push implements Channel { ... }
```

In OCaml a *column* is one function: `send` is one `match` holding the Email, Sms and Push cells; `format` the other three. In Java a *row* is one class: `Email` holds its own `send` and `format`. Adding WhatsApp adds a row: in OCaml two new cells, one in each existing function, both touched; in Java one new class, nothing else touched. Adding `schedule` adds a column: in OCaml one new function, nothing else touched; in Java one abstract method and a body in each of the three classes, all touched. Each way makes one direction free and the other expensive — that is the expression problem, and neither side is a mistake.

The point that is easily missed: in *both* languages the expensive direction is *checked*. Forget WhatsApp in OCaml's `format` and the compiler prints Warning 8 with the missing case named, once per function; forget `schedule` in Java's Push and the class does not compile until the method exists. A closed list of shapes lets the compiler enumerate the cells. A `switch` on a string, or an `if` chain on a kind field, has the same layout and none of the checking. Session 10 is about the ways to make both directions cheap at once.

Part 2 — at the keyboard: warmup.ml

The file `warmup/warmup.ml` holds five small types with holes marked `TODO`, and a check for each; the exercises are stated here, the file only says where to type. It needs no `dune`:

```
ocaml warmup/warmup.ml          # pass / FAIL / todo, one line per check
```

Each rung is a few lines, and each is one `match` on one type.

Exercise 5 (a variant)

Given the type of the days of the week, write the function that returns the next day (Sunday wraps to Monday) and the function that says whether a day is in the weekend:

```
type weekday = Mon | Tue | Wed | Thu | Fri | Sat | Sun

val next      : weekday -> weekday
val is_weekend : weekday -> bool
```

Write `next` with one arm per day and no wildcard, and `is_weekend` with an or-pattern. Then explain why seven arms are preferable to `| _ -> .`

Answer.

```
let next d = match d with
| Mon -> Tue
| Tue -> Wed
| Wed -> Thu
| Thu -> Fri
| Fri -> Sat
| Sat -> Sun
| Sun -> Mon

let is_weekend d = match d with
| Sat | Sun -> true
| Mon | Tue | Wed | Thu | Fri -> false
```

With seven arms, a day added to the type would make the compiler point here; a wildcard would accept it silently.

Exercise 6 (a record)

Given a point as a record with two fields, write the function that computes the distance between two points and the function that returns a point moved up by a given amount:

```
type point = { x : float; y : float }

val dist      : point -> point -> float
val move_up   : point -> float -> point
```

Write `move_up` with the copy-and-update form `{ p with y = ... }`. Then answer: after `move_up p 1.`, what is `p.y`, and why?

Answer.

```
let dist p q =  
  let dx = p.x -. q.x and dy = p.y -. q.y in  
  sqrt (dx *. dx +. dy *. dy)
```

```
let move_up p dy = { p with y = p.y +. dy }
```

`p.y` is unchanged: the functional update builds a new record; no field is `mutable`, so nothing could change the old one. Part C1 of the lab does exactly this to a pitch's octave.

Exercise 7 (payloads, and a constructor added)

Given a type of shapes whose constructors carry a radius, or a width and a height, write the two functions that compute a shape's area and its perimeter:

```
type shape =  
  | Circle of float  
  | Rect   of float * float  
  
val area      : shape -> float  
val perimeter : shape -> float
```

When both checks pass, **add** `| Square of float` to the type and run the file again without touching the functions. Copy down what the compiler prints. Then fix the functions, guided by what it printed.

Answer.

```
let area s = match s with  
  | Circle r    -> pi *. r *. r  
  | Rect (w, h) -> w *. h  
  | Square a    -> a *. a  
  
let perimeter s = match s with  
  | Circle r    -> 2. *. pi *. r  
  | Rect (w, h) -> 2. *. (w +. h)  
  | Square a    -> 4. *. a
```

Before the `Square` arms exist, the compiler prints, for `area` and again for `perimeter`: *Warning 8: this pattern-matching is not exhaustive. Here is an example of a case that is not matched: Square ..* Two sites, found without searching — exercise 4's grid in code. The lab's Part B2 does the same on six functions.

Exercise 8 (options)

Write two functions that may have no answer, and say so in their type: a division that returns `None` instead of failing on zero, and the largest element of a list, which the empty list does not have:

```
val safe_div : int -> int -> int option  
val max_of   : int list -> int option
```

In `max_of`, the recursive call returns an option: match on it, and treat `None` and `Some m` as the two cases they are.

Answer.

```
let safe_div a b = if b = 0 then None else Some (a / b)  
let rec max_of xs = match xs with  
  | [] -> None  
  | x :: rest -> (match max_of rest with  
    | None -> Some x  
    | Some m -> Some (if x > m then x else m))
```

'a option is “an 'a, or nothing”: two constructors, and the caller must match on them; there is no null value that a caller could forget to test.

Exercise 9 (a tree)

Given the recursive type of arithmetic expressions — a number, a sum of two expressions, or a product of two — write the function that evaluates an expression and the function that prints it with the fewest parentheses:

```
type expr =
  | Num of int
  | Add of expr * expr
  | Mul of expr * expr

val eval      : expr -> int
val to_string : expr -> string
```

Add (Num 1, Mul (Num 2, Num 3)) must print 1 + 2 * 3 and Mul (Add (Num 1, Num 2), Num 3) must print (1 + 2) * 3. When does a sum need parentheses? This tree — leaves and binary nodes — is the shape of the lab's melody, and the parentheses question returns in Part C2.

Answer.

```
let rec eval e = match e with
  | Num n      -> n
  | Add (a, b) -> eval a + eval b
  | Mul (a, b) -> eval a * eval b

let rec to_string e =
  let factor f = match f with
    | Add _ -> "(" ^ to_string f ^ ")"
    | _      -> to_string f
  in
  match e with
  | Num n -> string_of_int n
  | Add (a, b) -> to_string a ^ " + " ^ to_string b
  | Mul (a, b) -> factor a ^ " * " ^ factor b
```

A sum needs parentheses only inside a product, so the Mul arm looks at the *shape* of its children — a nested pattern — and wraps an Add. In the music language, Par is the node that must bring its own brackets.